

HPC101 – Hands-on Session

I. Navigating, Writing Jobs, Submission, and Monitoring

1. Download a repository as .zip file from <https://github.com/IITDhTraining/hpc101> and copy it to your home drive on the master node.
2. Log-in to your cluster.
3. Move the file copied from home drive to scratch space on your cluster.
4. Go to the directory where you have moved the file and extract the files from the .zip file that you have copied.
5. Go to the directory hpc101 (case-sensitive).
6. Write and submit a job script, `procinfo.job`, to find processor info (using the command `cat /proc/cpuinfo`) of the compute node on which you run the job. Your job should request 1 node and 1 core for a maximum time of 15 seconds. Rename the output file produced as `cpuinfo_compute.out`
 - a. How many physical cores are there in your compute node?
 - b. How many cores are there including hyperthreading?
7. Now find the processor info on the master node. Capture the output in a file called `cpuinfo_master.out`. Compare `cpuinfo_compute.out` and `cpuinfo_master.out` and note down your observations.
 - a. How many physical cores are there in your master node?
 - b. How many cores are there including hyperthreading?
8. Write and submit job script, `wordcount.job`, to sum the count of the number of words in each of the .txt files extracted (hint: use the commands `wc`, `cut`, and `bc`. Use pipes, backtick, and variables in shell). Your job should request 1 node and 1 core for a maximum time of 5 minutes.
 - a. Report the job status of the job submitted. When the job completes, what is the status that you see? rename the output file produced as `wordcount.out`
9. Write and submit a job script, `hostname.job`, to find the name of the compute nodes (using `hostname` command) on which you run the job. Your job should request 4 nodes and 1 core on each node. The job must run on all 4 nodes within a maximum time of 15 seconds. Rename the output file produced as `hostname.out`

II. Programming

1. Parallelize computing the value of π by numerical integration using the *Reimann sums* approach:

$$\text{Let } f(x) = \sqrt{1 - x^2} \text{ describe the quarter circle for } x = 0 \dots 1$$
$$\pi / 4 = \sum_{i=0}^{N-1} \Delta x f(x_i) \text{ where } x_i = i \Delta x \text{ and } \Delta x = 1/N$$

Sequential version of the program (`pi_seq.cpp`) is given to you. Write an OMP parallel program to do the same. Name your file as `pi_omp.cpp`

Background: You can compile your program as:
`g++ -fopenmp pi_omp.cpp -o pi_omp`

If the compilation is successful, you will see an executable file called `pi_omp`.
Execute the file using the command:

```
./pi_omp
```

2. Write an MPI parallel program to parallelize computing the value of pi using Reimann sums approach. You have been given `pi_mpi.cpp`. You need to fill in a line of code.
 - a. Use Reductions in MPI to implement your program. How many time steps would your implementation take to execute?
 - b. Run the program with 1, 2, 4, 8, 16, 32, and 64 processes on 1 node. Does the speedup you see in step match your expectations considering the processor configuration of your machine?

Background: MPI parallel programs can be written in C/C++/Fortran (and in Python as well). You compile a C/C++ program having MPI constructs as:

```
mpicc pi_mpi.c -o pi_mpi  
mpic++ pi_mpi.cpp -o pi_mpi
```

you run the program on a cluster of machines as follows.

If you are executing 10 copies of the executable `pi_mpi` you would simply run them as:

```
mpirun -np 10 pi_mpi
```