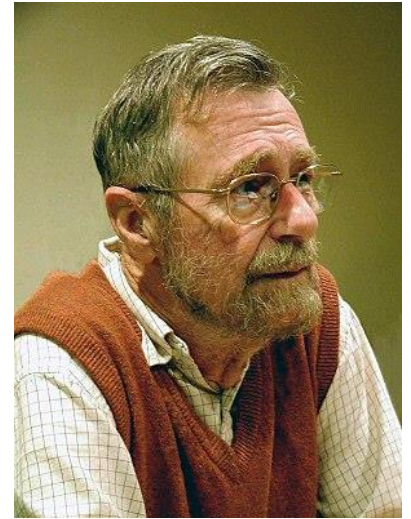"as long as there were no machines, programming was no problem at all;

when we had a few weak computers, programming became a mild problem, and

now we have gigantic computers, programming has become an equally gigantic problem."

- Edgar Dijkstra, 1972 Turing Award Lecture

A Gigantic Computer

- System 360 / Model 91

HPC Clusters and Programming

NSM Nodal Center for Training in HPC and AI, IIT Madras

Nikhil Hegde, IIT Dharwad

March 20, 2021

What is a Cluster ?

- Gigantic computer
  - from interconnecting several smaller computers

VIRGO Super Cluster, IIT Madras. Source: https://cc.iitm.ac.in/node/184

What is a Cluster ?

- Gigantic computer
  - from interconnecting several smaller computers
- Compute power in the order of 10^{15} floating point operations per second (Peta* FLOPS)
  - Your i7-based personal computer – few Giga FLOPs (10^9)

10^{12} is tera, 10^{18} is exa.
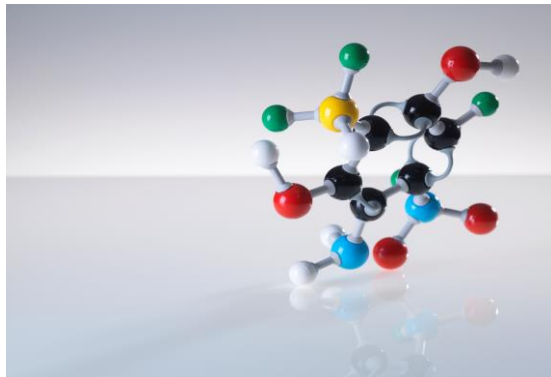
What is a Cluster ?

- Gigantic computer
  - from interconnecting several smaller computers
- Compute power in the order of 10^{15} floating point operations per second (Peta* FLOPS)
  - Your i7-based personal computer – few Giga FLOPs (10^9)
- E.g.
  - Chandra (IIT Palakkad), Virgo (IIT Madras), AnantGanak (IIT Dharwad) etc.

10^{12} is tera, 10^{18} is exa.
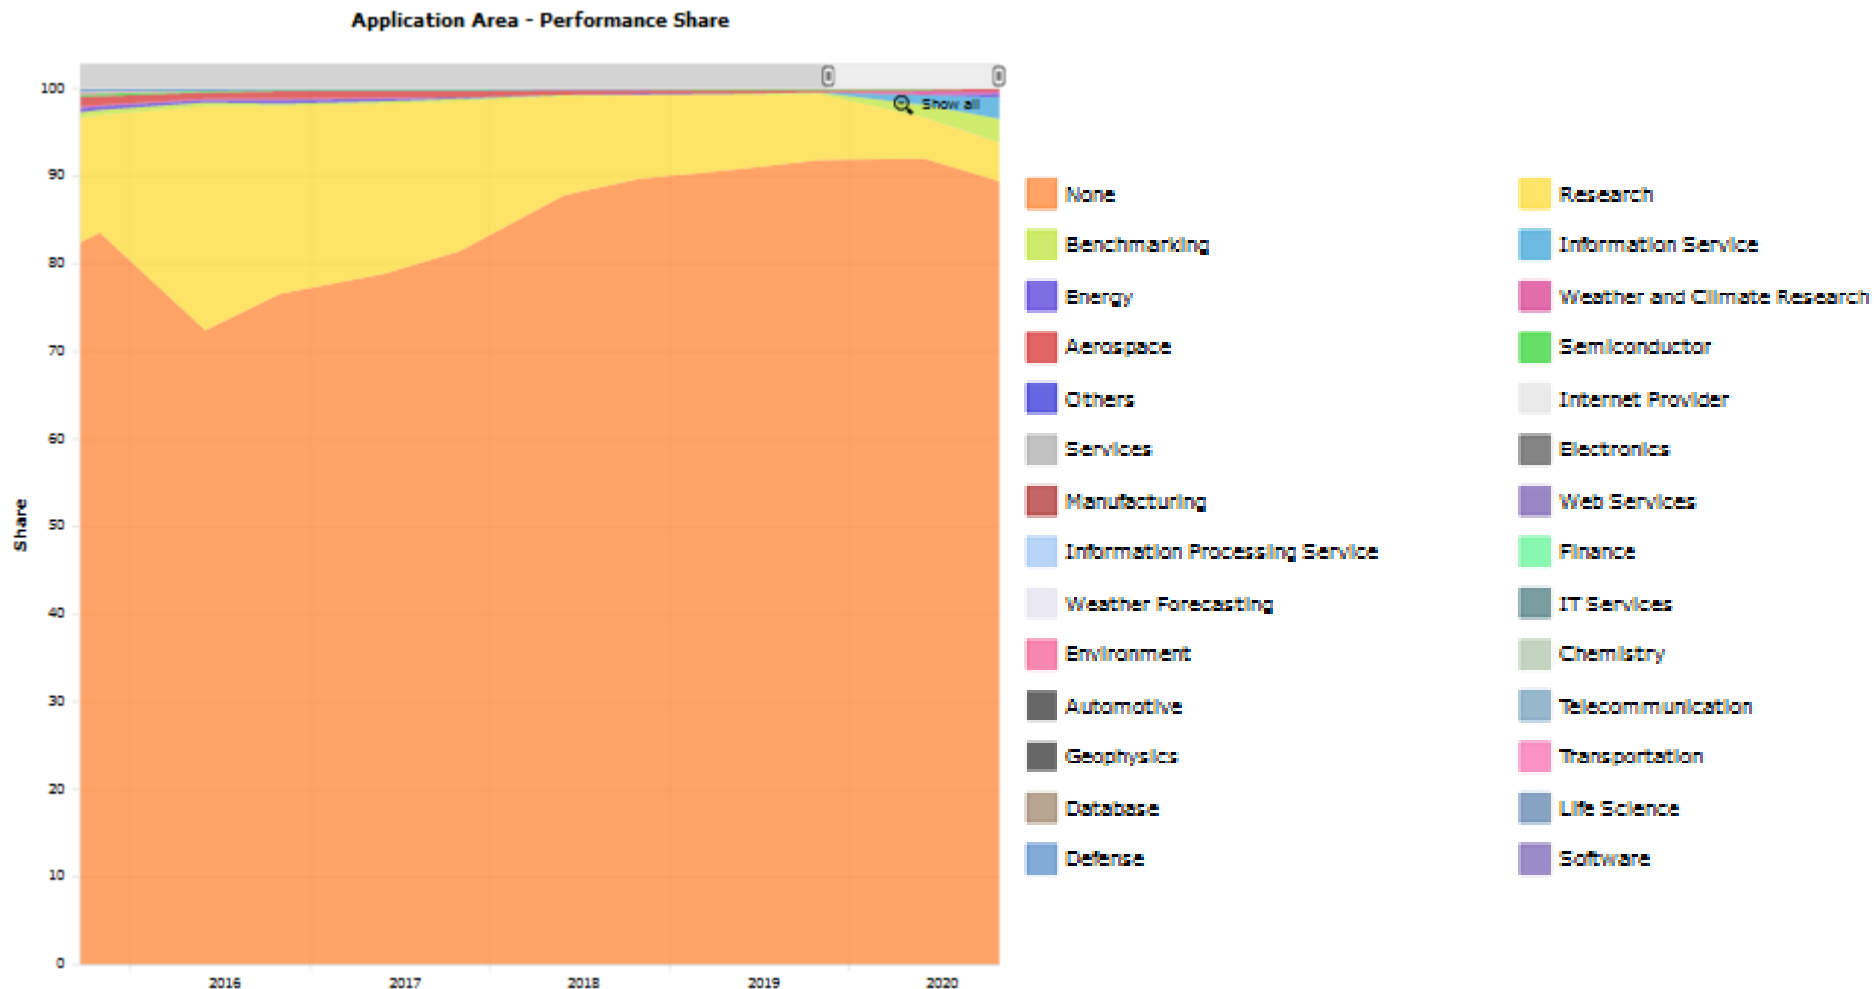
Why Clusters?

Financial Analysis

Design Simulation

Oil Exploration

Genomics

pic source: stock images

Weather Prediction

Why Clusters? Application Areas

Application Area - Performance Share

source: top500.org

Terminology - Cluster Elements

- Processing Element
  - Core, CPU, Node, GPU, Virtual CPUs
- Storage
- Interconnect

Hardware

- Partition
- Job and Job Scheduler
- Operating System (OS), Software Development Tools, Applications

Software

- *Infrastructure* – power, cooling,

Processing Element

- Node
  - Standalone computer
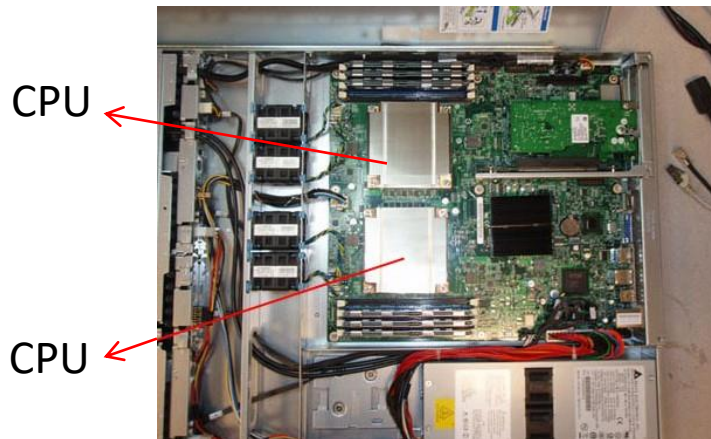 - Comprised of multiple CPUs/Processors/Cores, memory, network interfaces.

Each Green dot is a Node

  - `Master / Log-in` node: is what end user interacts with (think: operator's console)

Processing Element

- ## CPU/Processor and Socket
  - No consensus on terminology. Some vendors call multi-core CPUs as sockets.

  CPU →

  CPU →

  source: Blaise Barney, Introduction to Parallel Computing,
  https://hpc.llnl.gov/training/tutorials/introduction-parallel-computing-tutorial

  - Socket can also be a place to plug a CPU. E.g. dual-socket motherboard in pic.

Processing Element

- Core
  - Each PU shown is a core. Pic: IBM BG/Q with 18 Cores

source: Blaise Barney, Introduction to Parallel Computing,
https://hpc.llnl.gov/training/tutorials/introduction-parallel-computing-tutorial
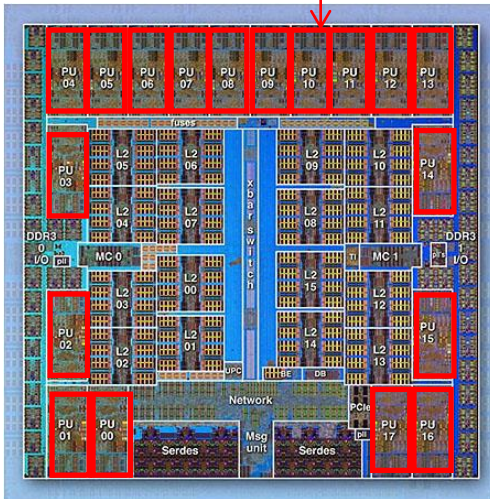
  - In the past, each CPU (with just one core) was a single execution unit. Now, each core is an independent execution unit.

Processing Element

- ## Thread (hardware)
  - Pathway for flow of instruction within a core
  - When exposed to the OS, the OS gives an illusion to the programmer that multiple cores exist ("HyperThreading")

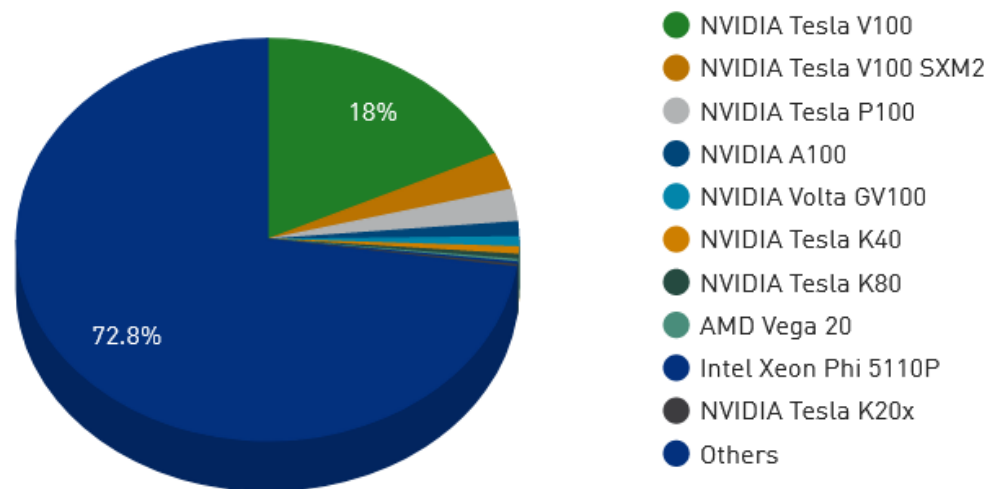Instruction sequence
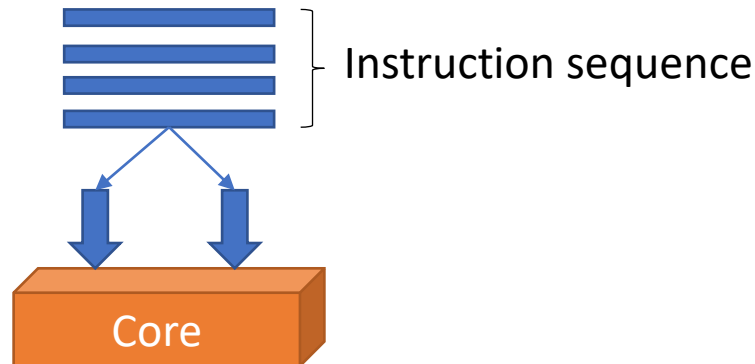
Core

Processing Element

- Virtual CPUs
  - A term that you often get to hear when working with 'Cloud' hosted systems
  - Virtual Machine (VM) assigned to a single physical core
    - A VM is an abstraction/emulation of a computer

Processing Element

- GPU (Graphics Processing Unit) – Add-on devices
  - Traditionally: accelerate image creation
  - Now: GPGPUs for large-scale modeling, genetic programming

Accelerator/Co-Processor System Share

- NVIDIA Tesla V100
- NVIDIA Tesla V100 SXM2
- NVIDIA Tesla P100
- NVIDIA A100
- NVIDIA Volta GV100
- NVIDIA Tesla K40
- NVIDIA Tesla K80
- AMD Vega 20
- Intel Xeon Phi 5110P
- NVIDIA Tesla K20x
- Others

source:top500.org

Cluster Elements

- Interconnect
  - Nvidia-Mellanox Switch

source: https://www.nvidia.com/en-in/networking/infiniband-switching/

  - Cabling (determines topology)
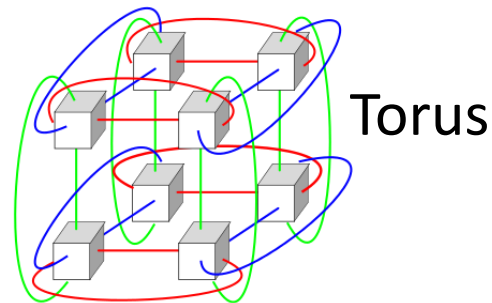
Torus

source: https://en.wikipedia.org/wiki/Torus_interconnect

https://www.rcac.purdue.edu/training/clusters101/

16

Cluster Elements

- Storage
  - $HOME
    - Landing directory when you log in to the master node
  - $SCRATCH and/or Parallel File System
    - A fast memory where you should keep all the data needed for executing the task
    - E.g. Lustre, BeeGFS, etc.

Cluster Elements - Software

- Job
  - A task performed by the cluster
  - Represented by a set of commands to the cluster, captured in a `script`, to precisely tell how to execute the task
  - Usually, the set of commands do not require your intervention i.e. non-interactive
    - You issue the commands (read: "submit a job") and go for coffee..

Cluster Elements - Software

- Batch System - Job Scheduler and Resource manager
    1. Provide a user interface to submit, monitor, and run jobs
    2. Manage the computational resources mentioned previously
    3. Implement the usage policies set by HPC Admin
- E.g. SLURM (Simple Linux Utility for Resource Management), PBS (Portable Batch System) – Torque, Moab.

Cluster Elements - Software

- Partitions/Queues
  - A logical grouping of (a subset of) nodes in the cluster
  - Single node can belong to multiple partitions (not done in practice)
  - Define attributes (and limits) for a job submitted to a Q
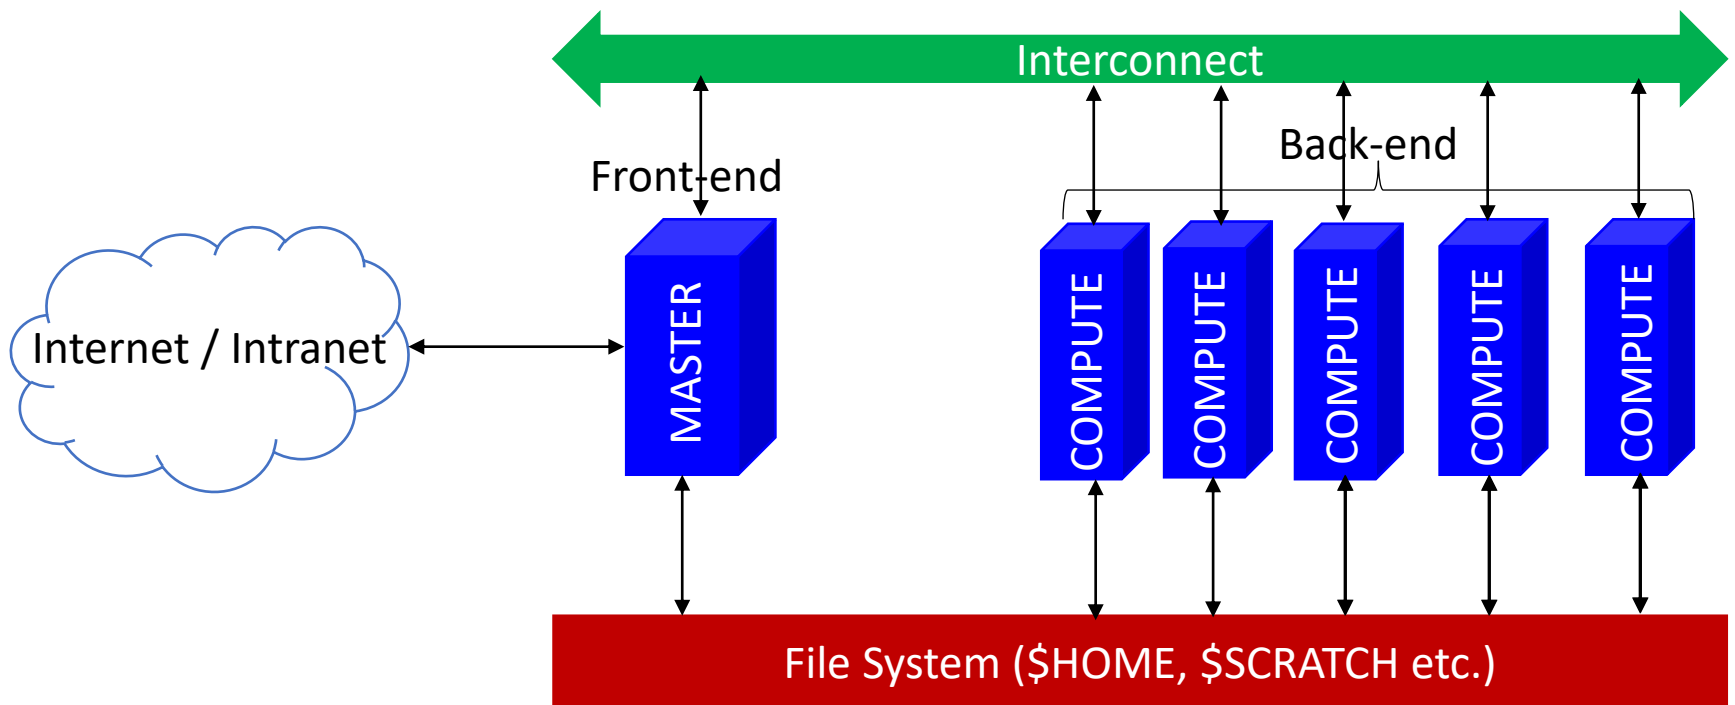   E.g.
    - A job submitted to a specific partition runs with low priority
    - A job can request a maximum of 4 cores
    - etc.

Cluster Elements - Software

- Operating System (OS), Software Development Tools, Applications
  - Linux-based OS in 100% of the supercomputing clusters in top500.org (2015 onwards)
  - Compiler tool chains, Runtime systems, Profilers
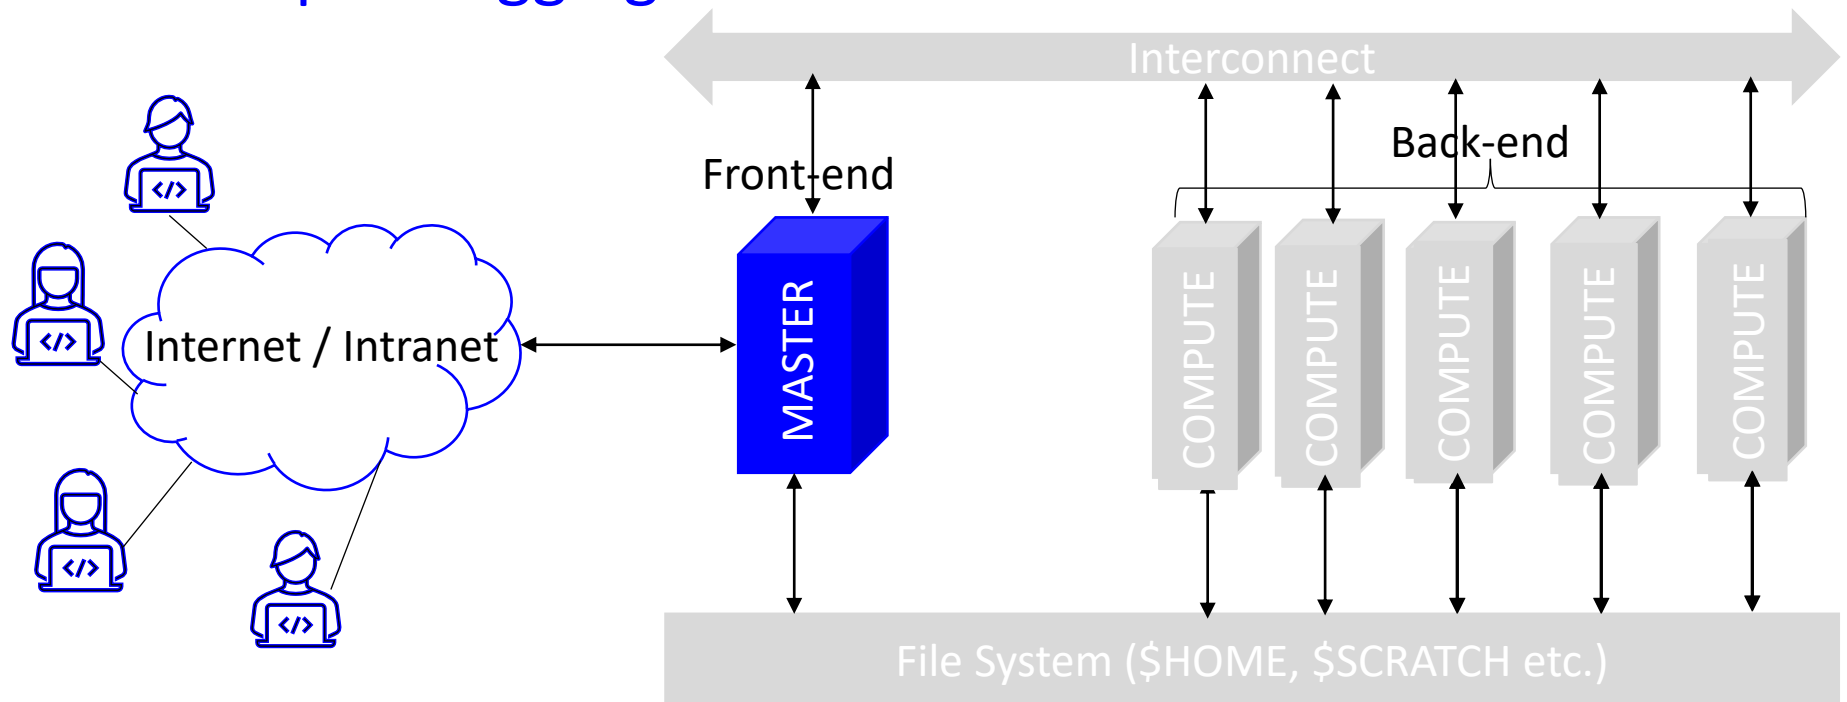 - E.g. GCC, ICC, MPICH, JDK, Docker, Matlab, Intel Parallel Studio, NVProf, Tau etc.

Recap

- Clusters

Clusters

- Step 1: Logging-in

Interconnect

Front-end

Back-end

MASTER

COMPUTE COMPUTE COMPUTE COMPUTE COMPUTE

Internet / Intranet

File System ($HOME, $SCRATCH etc.)

Logging-in

- Logging into remote Linux system (master node) requires you to use SSH ("Secure Shell Protocol")
  - Login credentials are encrypted
  - `SSH server` must be running on the system that you are logging into; Happens on most Linux systems by default.
  - `SSH client,` another piece of software, is used to authenticate and connect to the SSH `server`
    - Client software is available for all platforms (OSs)

Logging-in Windows

- Powershell on Windows 10
  - Press (Windows + R) -> Type "powershell"
  - Type "ssh <username>@<masternode_IP_address>
  - Type 'Yes' when prompted (only first time)
  - Provide log-in credentials

Windows PowerShell

```
Windows PowerShell
Copyright (C) Microsoft Corporation. All rights reserved.

Try the new cross-platform PowerShell https://aka.ms/pscore6

PS C:\Users\ndheg> ssh nikhilh@10.250.101.100_
```

Logging-in Windows

- PuTTY SSH client Windows
    - Download PuTTY from https://www.chiark.greenend.org.uk/~sgtatham/putty/latest.html (64-bit .exe)
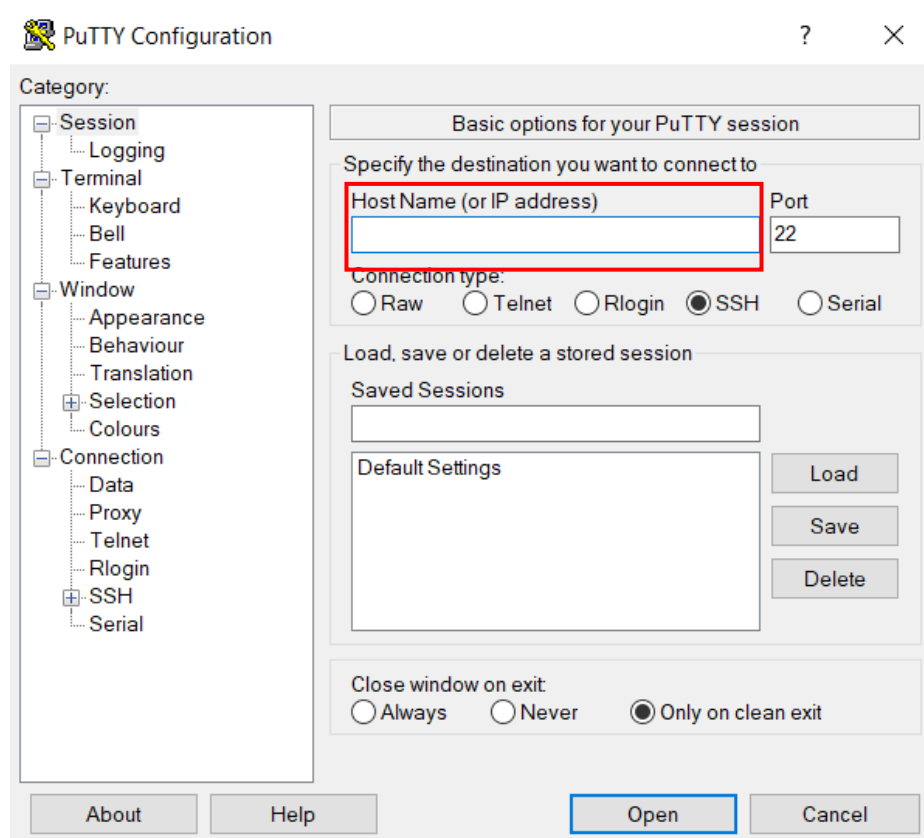   - Double click on the   icon after downloading
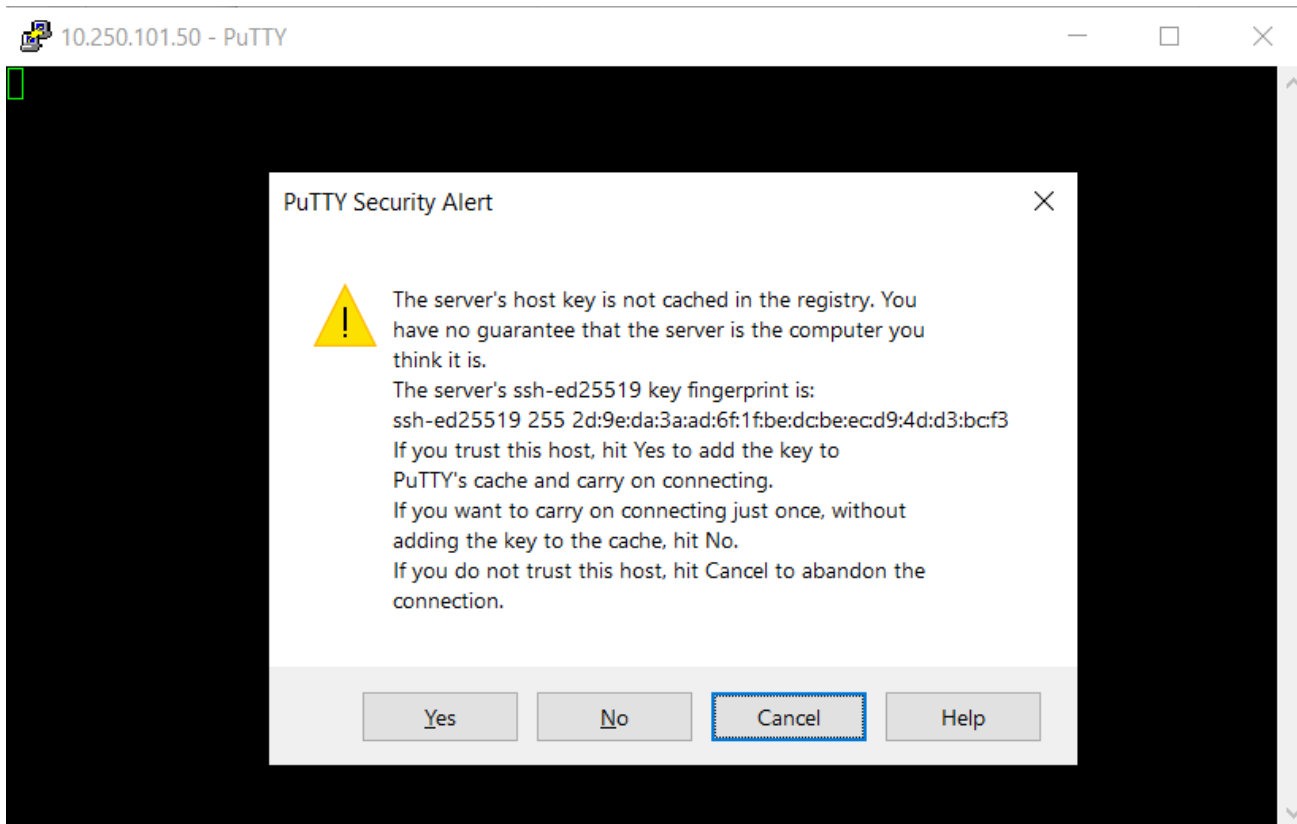
    putty

Logging-in Windows

- Type in the Host Name / IP address and click 'Open'
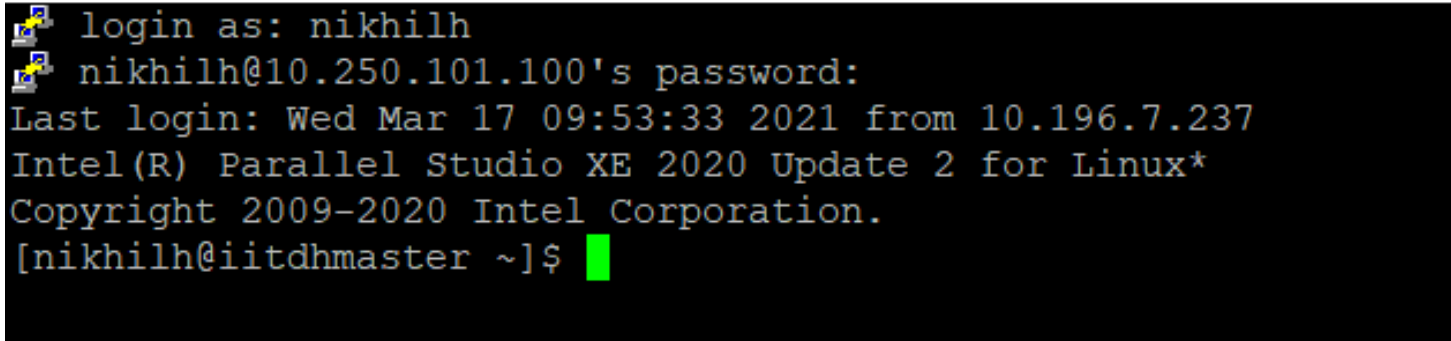
Logging-in Windows

- Click on 'Yes' (you are accepting the server host key)

Logging-in Windows

- Enter log-in credentials


```
nikhilh@iitdhmaster:~
  login as: nikhilh
  nikhilh@10.250.101.100's password:
Last login: Wed Mar 17 09:53:33 2021 from 10.196.7.237
Intel(R) Parallel Studio XE 2020 Update 2 for Linux*
Copyright 2009-2020 Intel Corporation.
[nikhilh@iitdhmaster ~]$
```

Logging-in MAC

- Open the 'Terminal' program on MAC (Go -> Applications -> Terminal)

```
Last login: Sun Mar  7 11:35:13 on ttys000

The default interactive shell is now zsh.
To update your account to use zsh, please run `chsh -s /bin/zsh`.
For more details, please visit https://support.apple.com/kb/HT208050.
apples-MacBook-Pro:~ apple$
```

- Type "ssh <username>@<masternode_IP_address>
- Type 'Yes' when prompted (only first time)
- Provide log-in credentials

Logging-in Linux

- If you are a Linux user, you know what a 'Terminal' is ☺

- Type "ssh <username>@<masternode_IP_address>

- Type 'Yes' when prompted (only first time)

- Provide log-in credentials

Clusters

- Step 2: Activities on the Master Node

Clusters

- **Step 2: Activities on the Master Node**

Interconnect

Front-end

Back-end

MASTER

COMPUTE COMPUTE COMPUTE COMPUTE COMPUTE

Internet / Intranet

File System ($HOME, $SCRATCH etc.)
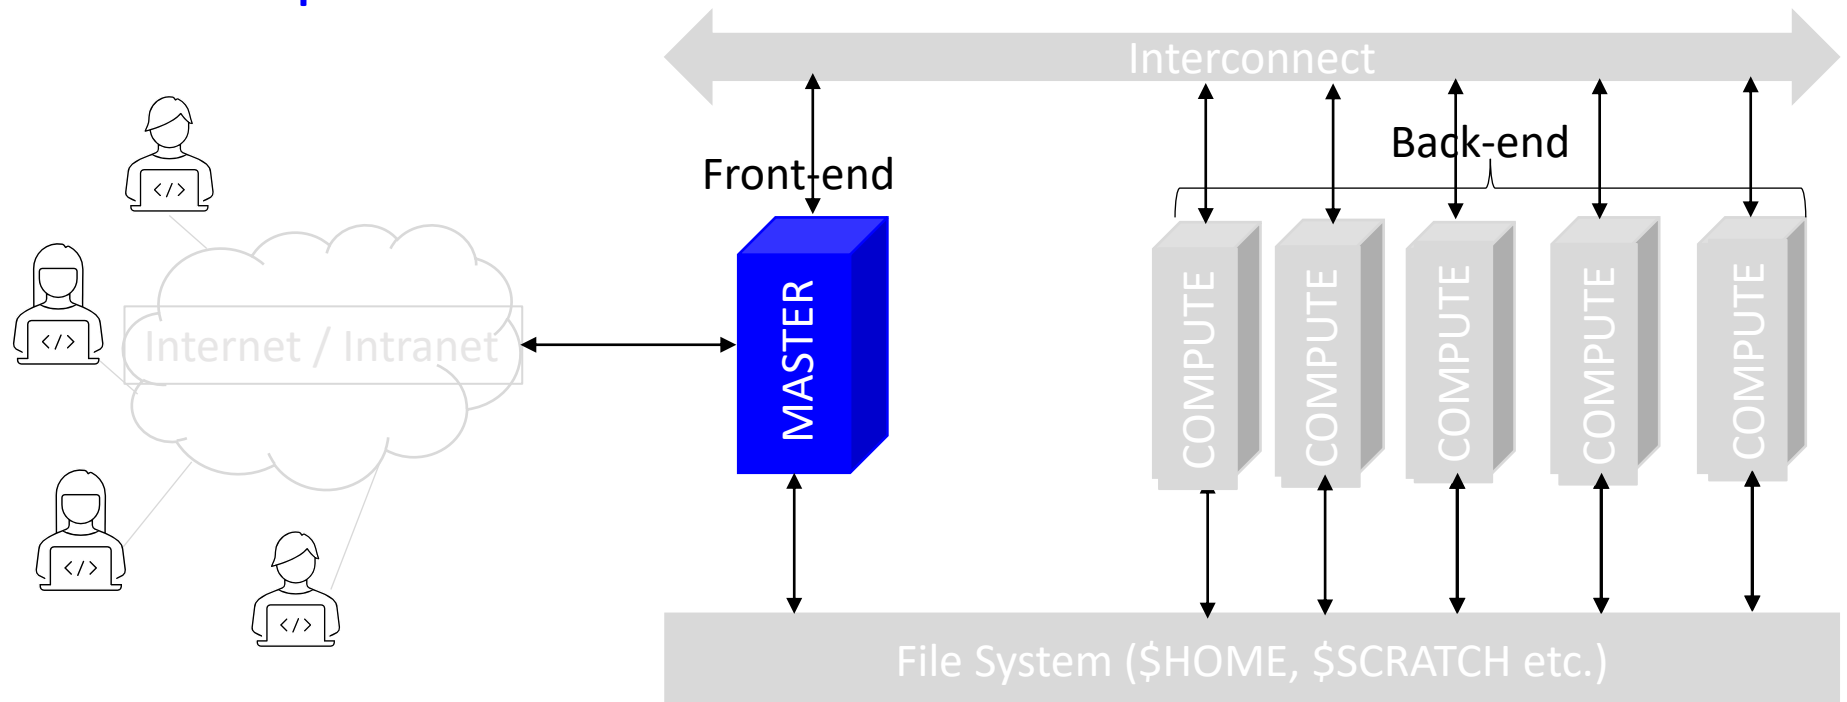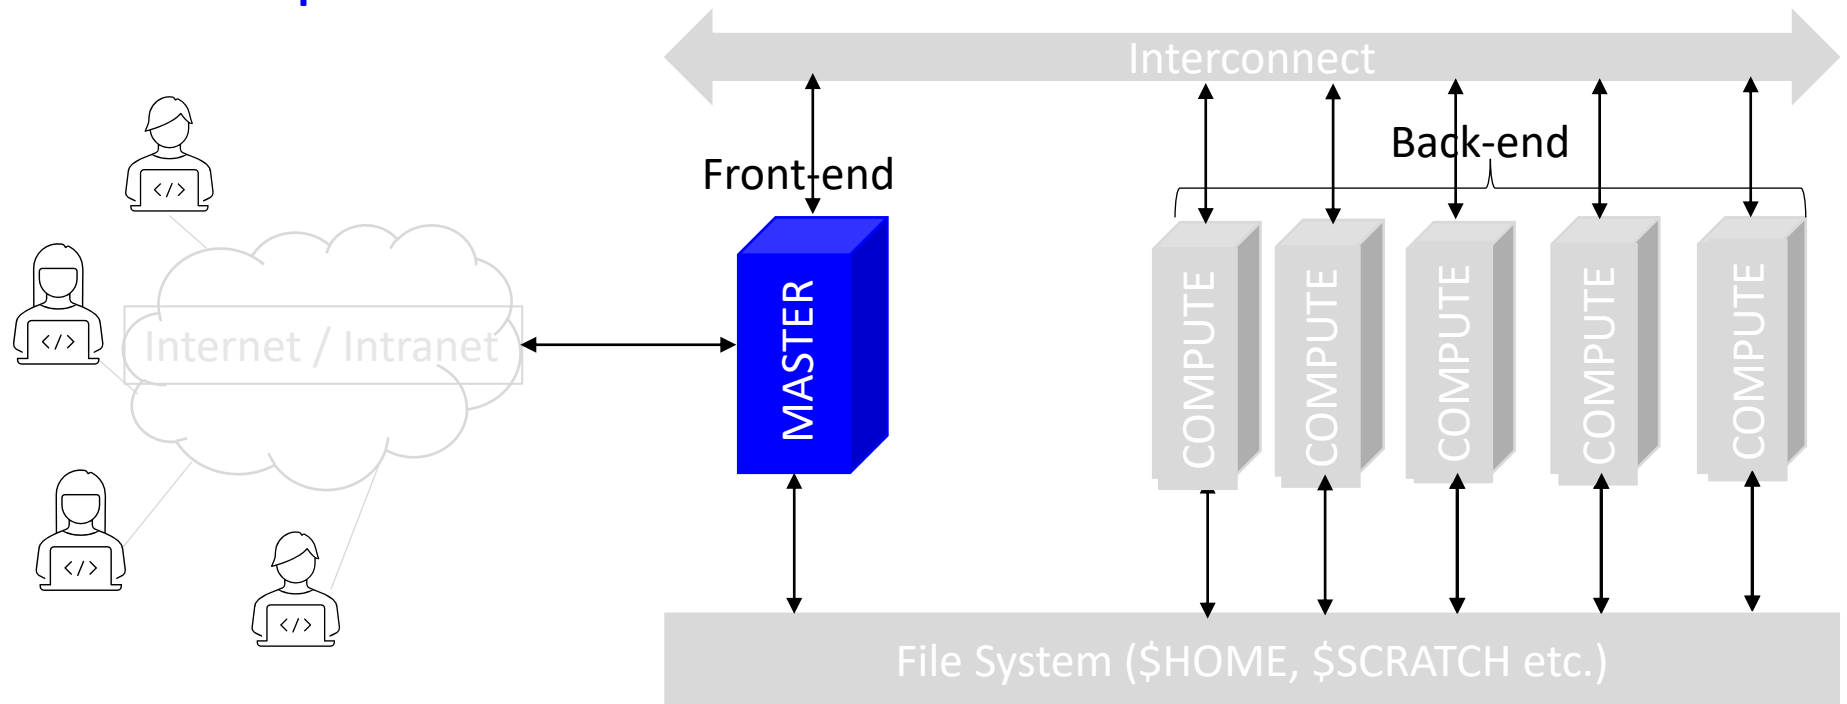
Master node runs Linux-OS. Let's review useful Linux commands

Useful Linux Commands

```
ls, ls -l
man
mkdir
cd
pwd
cp
mv
scp
rm  //use with caution!
```

```
cat
less
head, tail
vi, vim, emacs, nano, pico
gzip, tar, zip
who
```

Type "`man <command_name_here>`*" on the Linux terminal to get help info*

Useful Linux Commands

`scp` - To move files back and forth between master node and your local system

`scp file1 <user_name>@<master_node_ip>`

from your system to Master node

```
Windows PowerShell                                                    —  □  ×
PS C:\Temp> scp README.txt nikhilh@10.250.101.100:
nikhilh@10.250.101.100's password:
README.txt                                          100%  137    19.4KB/s  00:00
PS C:\Temp> scp nikhilh@10.250.101.100:pscp.exe .
nikhilh@10.250.101.100's password:
pscp.exe                                            100%  669KB   3.9MB/s  00:00
PS C:\Temp>
```

from Master node to your system

`scp <user_name>@<master_node_ip>:file1 .`

Useful Linux Commands

`zip, unzip` - To compress/uncompress folders/directories

`zip -r compressed.zip workshop_files/`

`unzip compressed.zip`

`-r` for recursively applying the compression to folders within

Useful Linux Commands

tar – **T**ape **Ar**chive to compress/uncompress folders/directories

`tar –cvf workshop.tar workshop_files/`

`tar -xvf workshop.tar`

Type `man tar` to know about flags

tar followed by `gzip` compression:

`tar –czvf workshop.tar.gz workshop_files/`

`tar -xzvf workshop.tar.gz`

Useful Linux Commands

Other utility commands

`cd HPC`   → change directory to HPC

`cd ..`   → change directory to parent

`vim hello.txt` → open a file hello.txt for editing. See vi commands

`ls` → list files in the current directory

`head hello.txt` → display the first few lines of hello.txt

`cat hello.txt` → display entire content of hello.txt

`pwd` → display the name of the present working directory

`who` → display the names of all users who have currently logged-in

Clusters

- Step 3: Job creation, execution, and monitoring
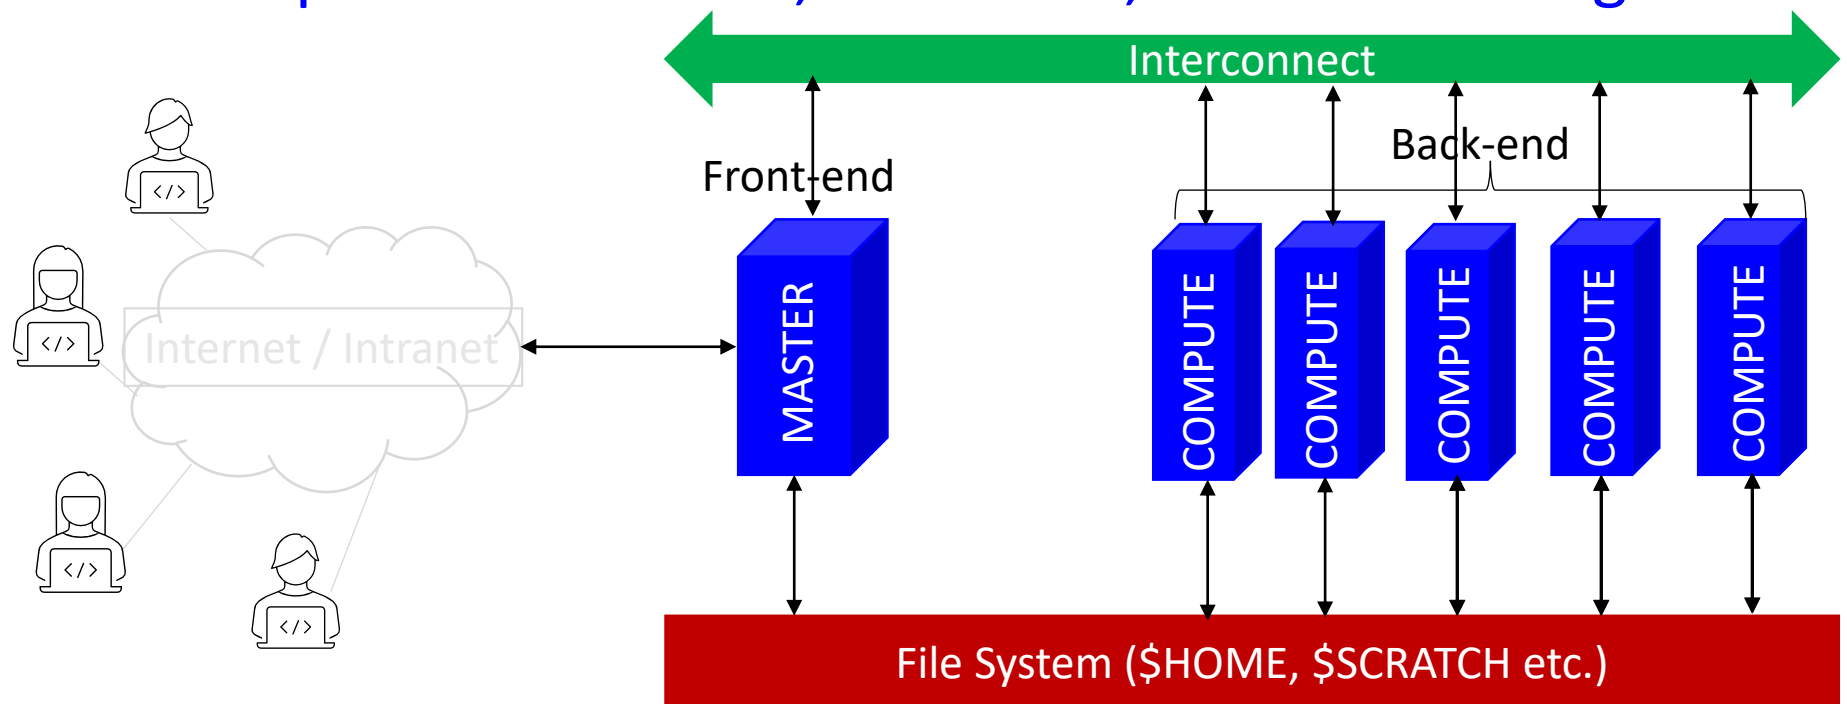
Clusters

- Step 3: Job creation, execution, and monitoring
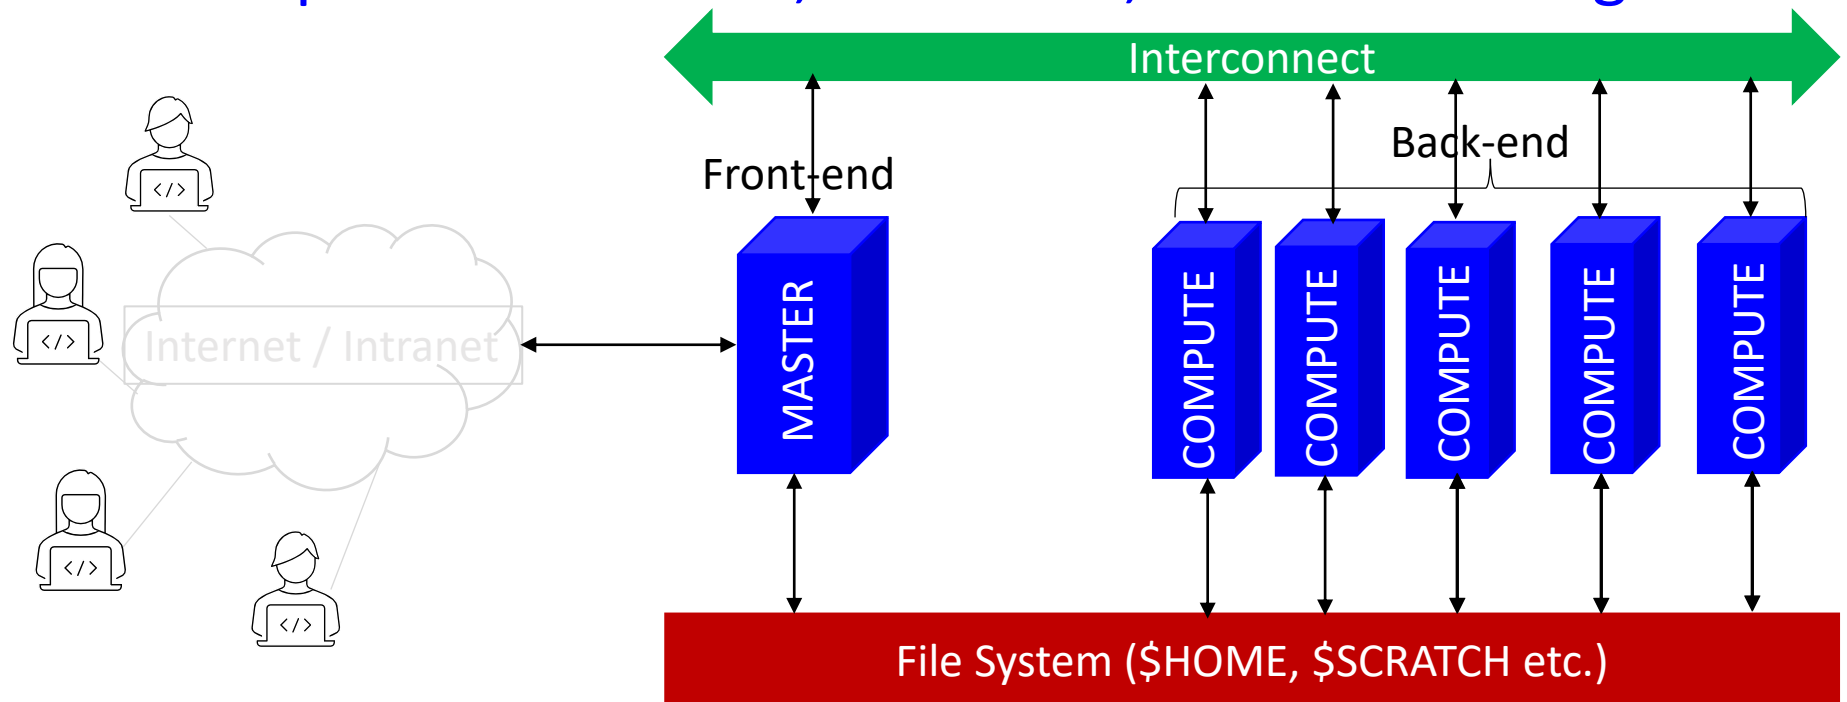
User must go through the Batch System for this step

Batch System (SLURM and PBS)

- Runs continuously (daemon) and monitors resources
- Scans the <u>queues</u> repeatedly
  - user jobs are submitted to queues
- Determines when and where the jobs are to be executed

MASTER NODE

Dispatches the job to one or more compute nodes

job

Is submitted to one of the Queues

Batch system

test

normal

long

Batch system

COMPUTE
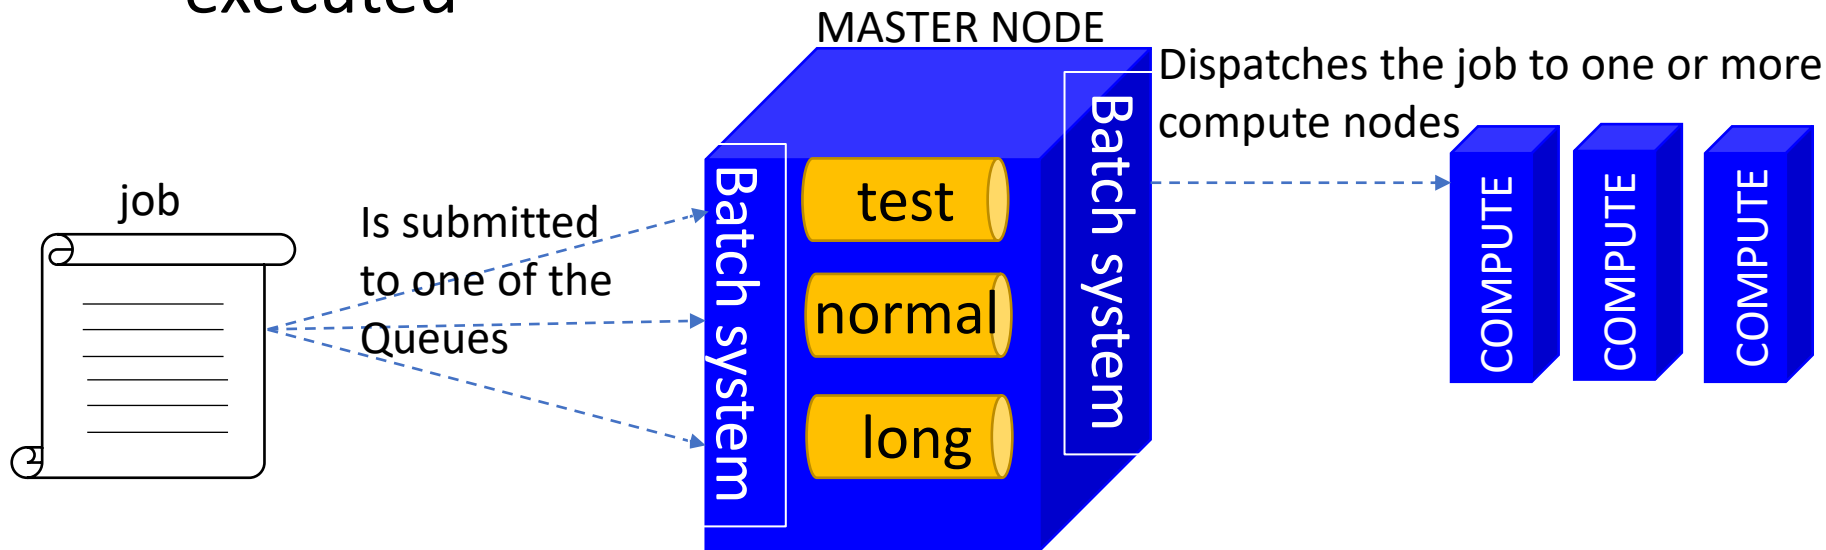
COMPUTE

COMPUTE

SLURM and PBS

- Provide exclusive or non-exclusive access to compute nodes
- Provide a set of user-commands for
  - job control
  - node control
  - accounting
  - Many other things …

SLURM Commands

- `sbatch` → to submit a job to the queue for later execution (batch mode)
- `scancel` → to cancel a queued or running job
- `sinfo` → report system info (queues, nodes, etc.)
- `squeue` → report job status
- `srun` → to create a job and execute
- `sattach` → to connect stdin/out/err to an existing job
- `sreport` → report resource usage by user, account etc.
- `sstat` → report accounting info about currently running jobs

SLURM Commands

- `salloc` → to create a job and start a shell to use it (interactive mode)

Life of a Job

1. Creation - job creation specifies:
   a. What resources are needed
   b. What commands will be run in the job
2. Submission - Successful submission implies:
   - Creating a job object/entry in the queue
   - Returning an ID associated with the object
3. Execution - when requested resources become available:
   - Job is launched on compute nodes
   - The commands specified in 1.b above are executed in sequence
4. Completion
   - Can happen because of normal completion of execution, abnormal termination, or job exceeding requested resources
   - output (stdout/err) stored in file at job submit location by default

Hands-on Session

- visit https://github.com/IITDhTraining/Cluster101 and download the repository as .zip file to get started
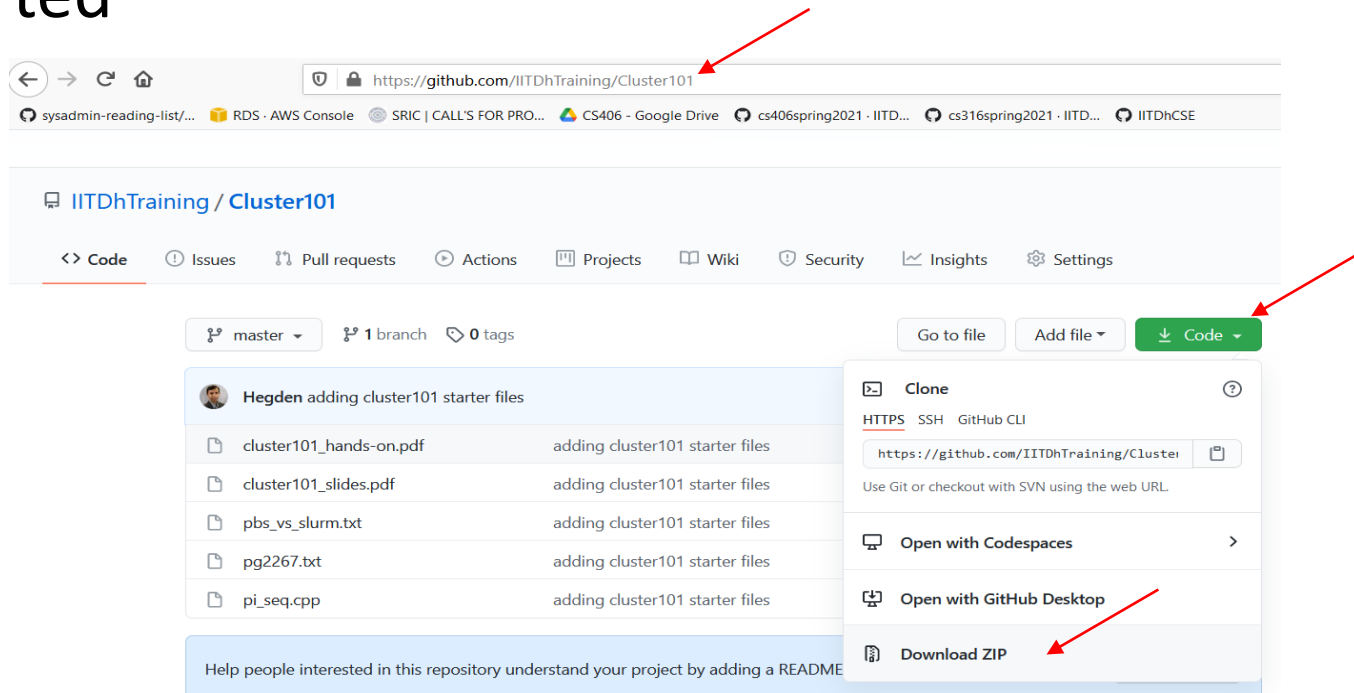
Clusters – How are they programmed?

- Programming for clusters is (all) about exploiting massive amount of hardware parallelism!

- Program needs to be executed

- Processes and Threads are units of execution

Threads and Processes

- Abstraction provided by OS
- Process
  - Self-contained i.e. has its own private resources to execute/run programs. E.g. of a resource: memory.
  - Is an instance of a running program.
  - Have an illusion that *entire computer* is for itself.
- Thread
  - Belongs to a process. Share memory and other resources among threads of the same process.
  - Have an illusion that *entire processor* is for itself.

Clusters – How are they programmed?

- It's (all) about parallelism!
- Sequential Program - *single sequence of instructions - single-threaded.*

source: Blaise Barney, Introduction to Parallel Computing,
https://hpc.llnl.gov/training/tutorials/introduction-parallel-computing-tutorial
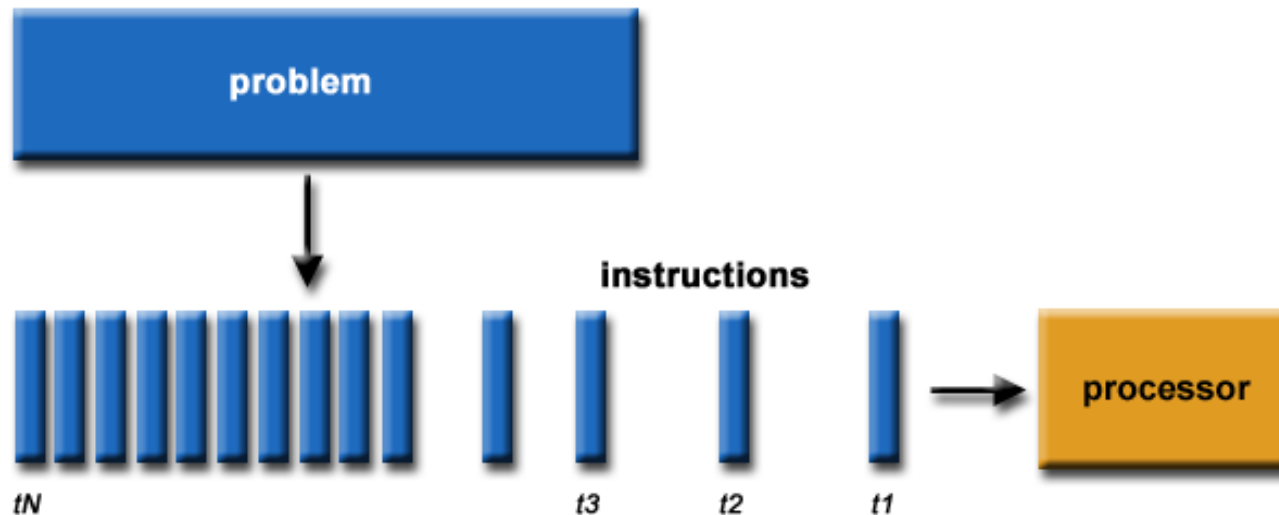
Clusters – How are they programmed?

- It's (all) about parallelism!

- Sequential Program - *single sequence of instructions - single-threaded.*

- Concurrent Program - *Multiple sequence of instructions executing concurrently.* Instructions from one sequence *may* communicate and interfere with other.
  - How are they (multi- threads) executing?
    - Multiprogramming – threads multiplexing their execution on **one processor**
    - Multiprocessing – threads multiplexing their execution on multiprocessor or **multicores**
    - Distributed Processing – processes multiplexing their executions on **multiple nodes**

Clusters – How are they programmed?

- Parallel Program – a *concurrent program* designed to execute on **parallel hardware**
  - Multiple processors in a computer (multiprocessing),
  - Multiple computers in a network (distributed processing)
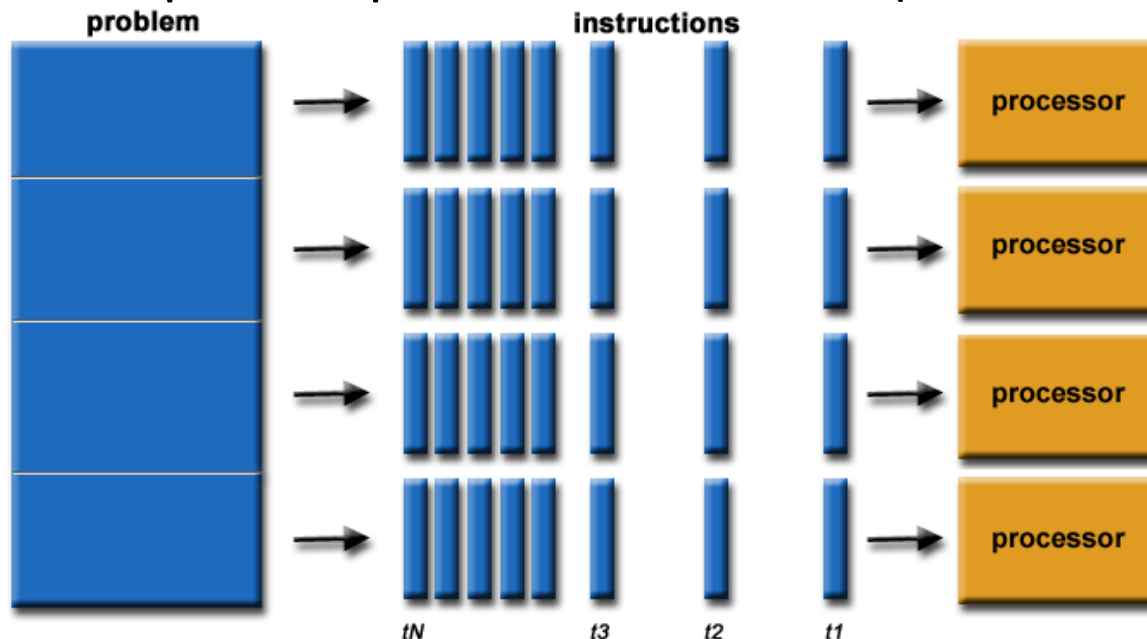
source: Blaise Barney, Introduction to Parallel Computing,
https://hpc.llnl.gov/training/tutorials/introduction-parallel-computing-tutorial

Parallel Hardware

- Flynn's taxonomy - categories of computing systems
    - Based on how processing elements see *instruction* and *data*

	Single Data (SD)	Multiple Data (MD)
Single Instruction (SI)	SISD	SIMD
Multiple Instruction (MI)	MISD	MIMD

Parallel Hardware

- Flynn's taxonomy - categories of computing systems
    - Based on how processing elements see *instruction* and *data*

	Single Data (SD)	Multiple Data (MD)
Single Instruction (SI)	SISD	SIMD
Multiple Instruction (MI)	MISD	MIMD

Clusters belong to this category

Parallel Hardware

- Categorization based on how processing elements see *system memory*
  - Shared Memory
  - Distributed Memory
  - Distributed-Shared Memory

Parallel Hardware

- Categorization based on how processing elements see *system memory*
  - Shared Memory
  - Distributed Memory
  - Distributed-Shared Memory

Most clusters belong to this category

Parallel Hardware

Processor 1	Processor 2

Memory

Processor 4	Processor 3

⬅ **Shared Memory Architecture**

`←interconnect`

Processor 1	Processor 2
Memory	Memory
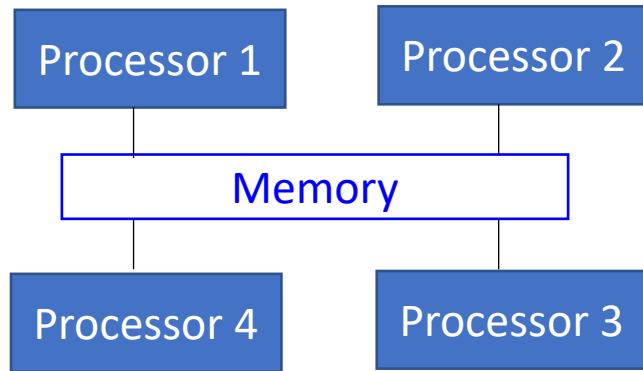
Distributed Memory Architecture ➡

Processor 4	Processor 3
Memory	Memory
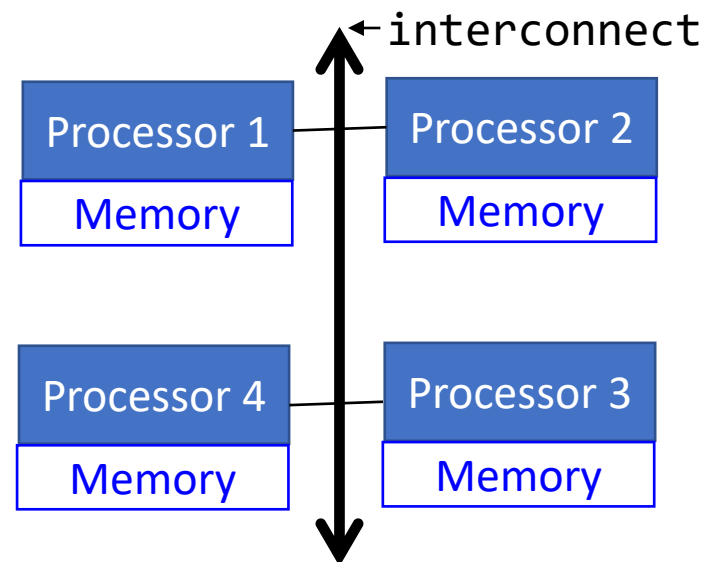
Data and Control Parallelism

- Threads executing the same function but with different data – *data parallelism*
  - E.g. two construction workers laying bricks to build walls of different parts of a house

- Threads executing different functions – *control parallelism*.
  - E.g. A carpenter getting a window frame ready while a mason is laying bricks in the wall

Need for Open MP
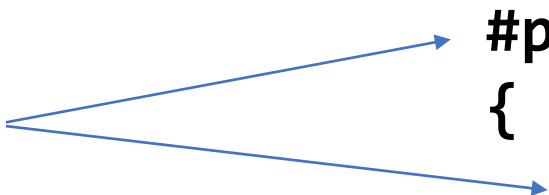
- Multithreaded programs using `std::thread` or `pthreads` :
  - do not *scale automatically* - add more processors and you will have to rewrite the program to utilize available processors
  - (pthreads) programs are not portable

- Open MP (multiprocessing) / OMP provides a scalable and portable alternative to data-parallel computing on *shared-memory architectures*

What is OpenMP

- An open standard for shared memory programming in C/C++ and Fortran

- Supported by IBM, Intel, GNU and others

- Same program running on multiple threads each operating on different data (single program multiple data – SPMD)

Programming in OpenMP

Compiler directives

```
#pragma omp parallel
{
    #pragma omp for
    for(i=0;i<N;i++)
        fruits[i].Energy();
}
```

- `#pragma parallel`

  executes as many threads as there are processors

- `#pragma omp for`

  divides the whole work among available threads

Programming in OpenMP

```
#pragma omp parallel
{
    #pragma omp for
    for(i=0;i<N;i++)
        fruits[i].Energy();
}
```

- Example of loop parallelism
  - Common in scientific codes
- Programmer is still responsible for handling data races.

Programming in OpenMP

```
//code region 0
#pragma omp parallel
{
    //code region 1
    #pragma omp for
    for(i=0;i<N;i++) {
        //code region 2
    }
    //code region 3
}
//code region 4
```

Master thread

Thread Number:1 2 3 4 5 . . . P   Workers

Workers

Workers

Master thread

- Synchronization point
/ **barrier**
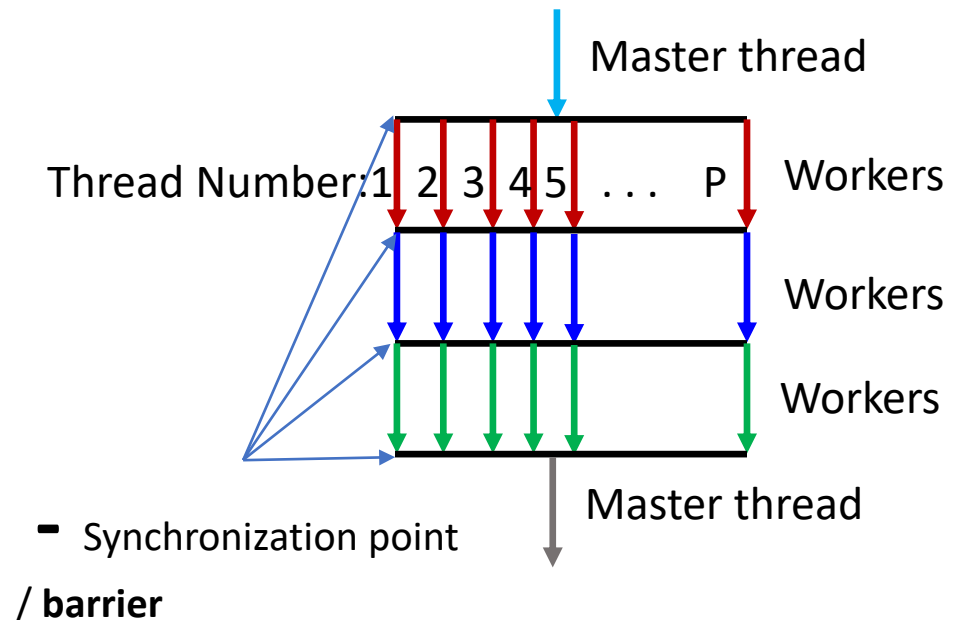
- Execution begins with a master thread (executes code region 0 and 4)
- Master thread creates / forks worker threads (execute code region 1,2, and 3)
- Worker threads join master thread

fork / join parallelism

62

Programming in OpenMP

- How many workers? / threads?
  - = number of processors by default. Can also be set with `omp_set_num_threads(P)`
  - Can query the number of processors available on a machine with `omp_get_num_procs()`
  - Each thread has an ID returned by `omp_get_thread_num()`

Programming in OpenMP

- Example – *reductions (sum of array elements)*

```c
#include<omp.h>
int main() {
    int total=0, a[NUM_SAMPLES];
    for(int i=0;i<NUM_SAMPLES;i++)
        a[i] = i;
    #pragma omp parallel
    {
        //#pragma omp for reduction(+: total)
        for(i=0;i<NUM_SAMPLES;i++)
            total = total + a[i];
    }
}
```

Programming in OpenMP

- Other operations supported in reductions:
  - +: addition
  - *: multiplication
  - |: bitwise OR
  - &: bitwise AND
  - ^: bitwise exclusive OR
  - ||: logical OR
  - &&: logical AND

  Note the *commutative nature* of these operations

Programming in OpenMP - Summary

- Open MP provides a way to specify what parts of program execute in parallel with one another

- How the work is distributed across different cores

- Whether to serialize (atomic) accesses to memory

- What order memory is read and written (barriers – nowait clause)

   All while providing *portable performance*

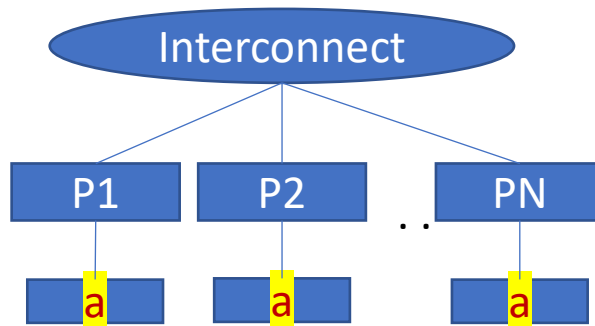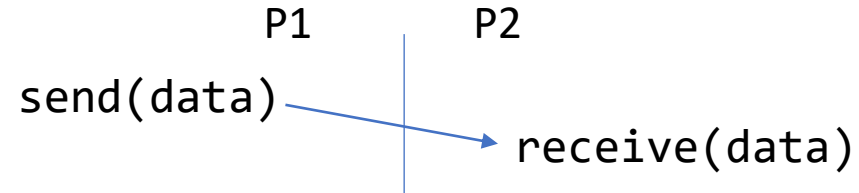# Distributed Memory Programming

- Program executes as a collection of processes

  Distributed Processing – processes multiplex their executions on multiple machines

- Each process / processor has its own memory

  - Total memory available for an MPI program is the combined memory space of all processors

  - Exchanging data requires cooperation between two processors

Distributed Memory Programming

- Data exchange requires explicit communication:

P1   P2

send(data)

receive(data)

Programmer must set up communication channels and exchange data

Interconnect

P1   P2   ..   PN

a   a   a

Value of a in P1 may be different from that in P2

1) P1 sends a copy of a to P2,

2) P2 receives the copy stores it in its data region.

- Every data element must belong to one of the memory spaces Programmer must decide where to place data

Distributed Memory Programming

- More complex than shared-memory programming model

- Most programs are written in *Single Program Multiple Data* (SPMD) model

- Computing power and cost scaling is better than with shared memory – e.g. rack mounted blades

- E.g. weather forecasting, simulating dynamics of gases and fluids - where to put exhaust fans in a basement parking? can an ATV topple while wading through body of water?

Distributed Memory Programming - Example
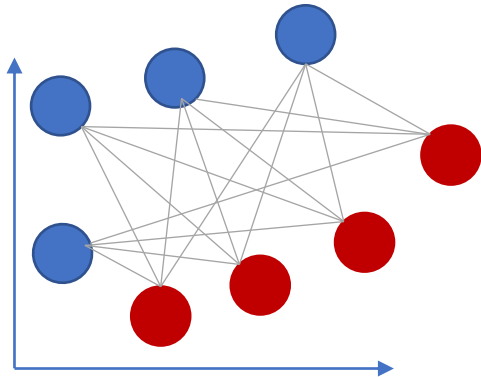
Calculate the distance from each point in set Blue (B) to that in set Red (R) and store the result in set D

```
for(i=1 to 4)
    for(j=1 to 4)
        D[i][j] = distance(B[i],R[j])
```

P1

Send(B[3..4], R[1..4], P2)

D   B

D

R

Receive(D[3..4][1..4], P2)

R   **B**

(is divided)

Aggregate

P2

Receive(B[3..4], R[1..4], P1)

D   B

R

Send(D[3..4][1..4], P1)

compute
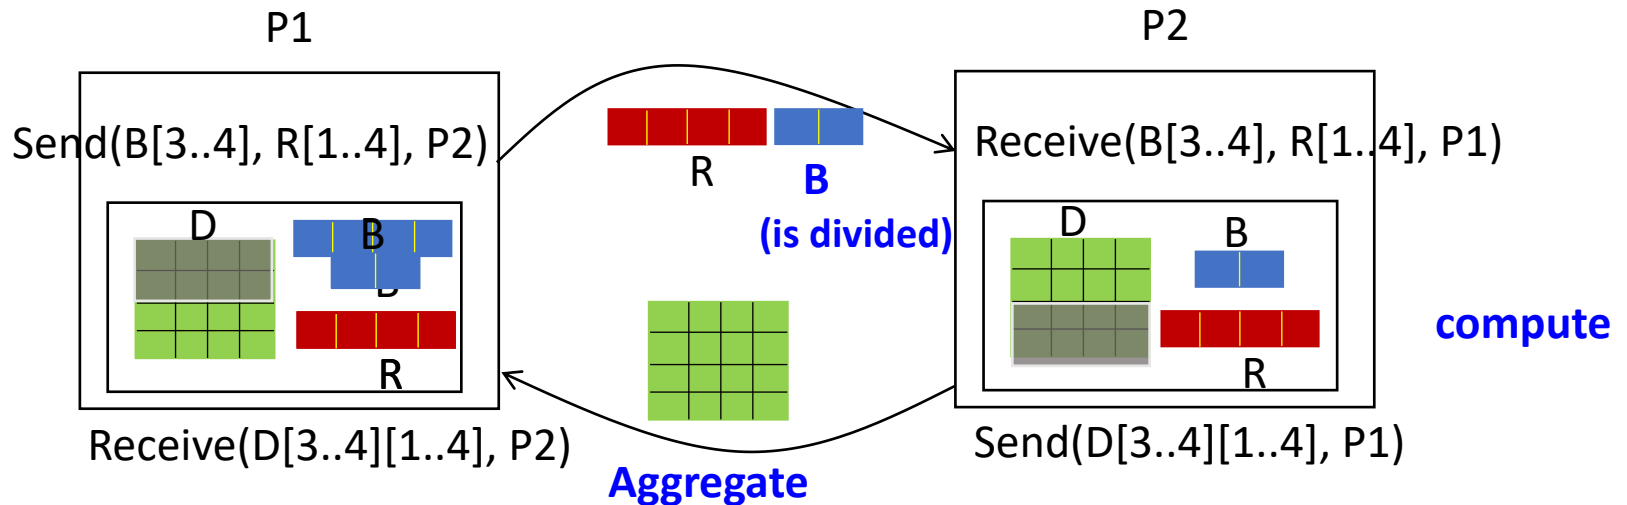
Distributed Memory Programming - Example

```
        Processor P1
//initialize B, R, and D
Send(B[3..4], R[1..4], P2)
for(i=1 to 2)
    for(j=1 to 3)
        D[i][j] = distance(B[i], R[j])
Receive(D[3..4][1..4], P2)
```

```
        Processor P2
//initialize B, R, and D
Receive(B[3..4], R[1..4], P1)
for(i=3 to 4)
    for(j=1 to 3)
        D[i][j] = distance(B[i], R[j])
Send(D[3..4][1..4], P1)
```

- How is work divided among processors?
- What does it mean for send and receive to complete?
- How does a receiver interpret data that a sender sends?

Converting to MPI Program

- Initialization and Termination

```
#include<mpi.h>

int main(int argc, char* argv[]) {
    MPI_Init(&argc,&argv); //initializes MPI Environment

    //all other code here

    MPI_Finalize(); //releases system resources
}
```

Converting to MPI Program

- Environment
  - 0,1,. . 5 – ranks / process numbers
  - MPI_COMM_WORLD – communicator / group of processes that are allowed to exchange messages
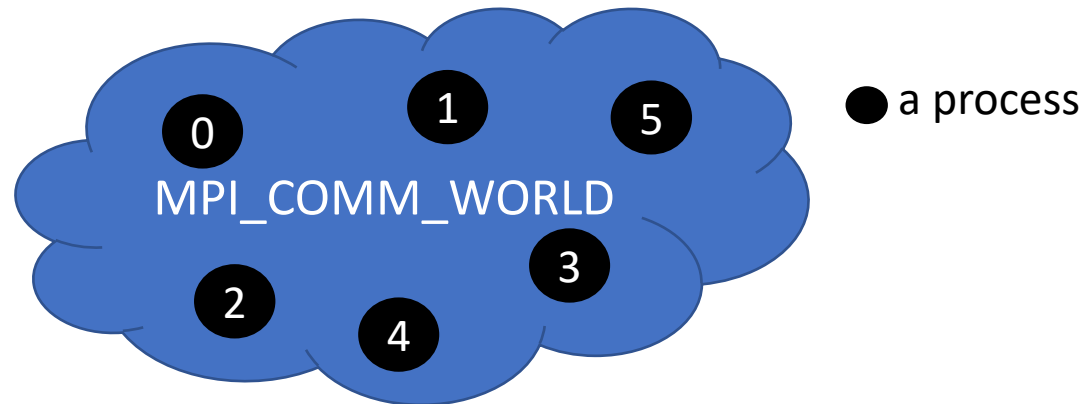
a process

- MPI_Init initializes the communicator.

Converting to MPI Program

- Rank and Size: obtaining process number and number of processes in the execution environment

```
#include<mpi.h>
int main(int argc, char* argv[]) {
    int rank, size;
    MPI_Init(&argc,&argv);
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Finalize();
}
```

Converting to MPI Program

- Divide the work and compute

```
#include<mpi.h>
int main(int argc, char* argv[]) {
    int rank, size;
    float d[NUMPOINTS][NUMPOINTS], b[NUMPOINTS], r[NUMPOINTS];
    //initialize b and r arrays from input
    MPI_Init(&argc,&argv);
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    for(i=rank*NUMPOINTS/size; i<(rank*NUMPOINTS/size)+NUMPOINTS/size; i++)
        for(j=0;j<NUMPOINTS;j++)
            d[i][j] = distance(b[i], r[j]);
    MPI_Finalize();
}
```

Converting to MPI Program

- Aggregate the results at master – MPI_Send, MPI_Recv

```
#include<mpi.h>
int main(int argc, char* argv[]) {
    int rank, size;
    float d[NUMROWS][NUMROWS], b[NUMROWS], r[NUMROWS];
    //initialize b and r arrays from input
    MPI_Init(&argc,&argv);
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    //compute d[i][j] as before

    if(rank !=0){
     for(int i=rank*NUMPOINTS/size;i<(rank*NUMPOINTS/size + NUMPOINTS/size);i++)
        MPI_Send(d[i], NUMPOINTS, MPI_FLOAT, 0, MY_MESSAGE_TAG, MPI_COMM_WORLD);
    else
     for(int i=NUMPOINTS/size;i<NUMPOINTS;i++)
        MPI_Recv(d[i], NUMPOINTS, MPI_FLOAT, MPI_ANY_SOURCE, MPI_ANY_TAG,
MPI_COMM_WORLD, &status);
    MPI_Finalize();
}
```

Converting to MPI Program

- Aggregate the results at master – MPI_Send, MPI_Recv

```
#include<mpi.h>
int main(int argc, char* argv[]) {
    int rank, size;
    float d[NUMROWS][NUMROWS], b[NUMROWS], r[NUMROWS];
    //initialize b and r arrays from input
    MPI_Init(&argc,&argv);
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    //compute d[i][j] as before

    if(rank !=0){
      for(int i=rank*NUMPOINTS/size;i<(rank*NUMPOINTS/size + NUMPOINTS/size);i++)
        MPI_Send(d[i], NUMPOINTS, MPI_FLOAT, 0, MY_MESSAGE_TAG, MPI_COMM_WORLD);
    else
      for(int i=NUMPOINTS/size;i<NUMPOINTS;i++)
        MPI_Recv(d[i], NUMPOINTS, MPI_FLOAT, MPI_ANY_SOURCE, MPI_ANY_TAG,
MPI_COMM_WORLD, &status);
    MPI_Finalize();
}
```

Send buffer

Count of sent items

Data type of sent items

Destination process ID

Message ID

Recv buffer

Recv count

any Message ID

From any Source Process

Converting to MPI Program

- Aggregate the results at master (collectives) – MPI_Gather

```
#include<mpi.h>
int main(int argc, char* argv[]) {
    int rank, size;
    float d[NUMROWS][NUMROWS], tmpd[NUMROWS], b[NUMROWS], r[NUMROWS];
    //initialize b and r arrays from input
    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    //compute distance into tmpd[NUMROWS];

    MPI_Gather(tmpd, NUMROWS, MPI_FLOAT, d, NUMROWS, MPI_FLOAT, 0,
MPI_COMM_WORLD)
    MPI_Finalize();
}
```

Send buffer   Count of sent items   Data type of sent items   Source Process ID
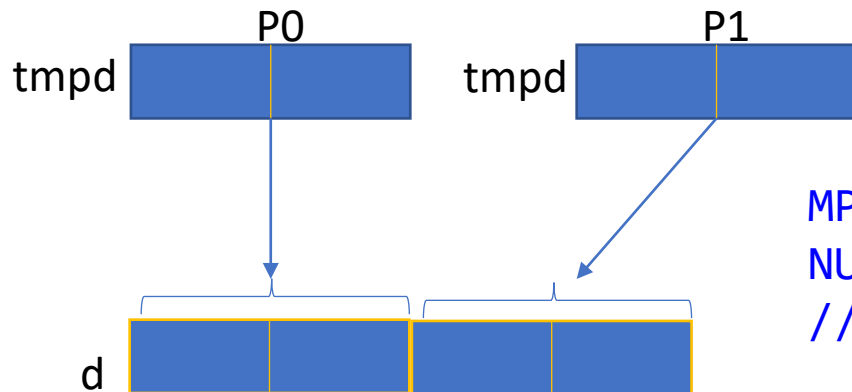
Recv buffer   Recv count

Converting to MPI Program

- collectives – MPI_Gather

MPI_Gather(sendbuf, sendcnt, send_type, recvbuf, recvcnt, recv_type, source_proc, MPI_COMM_WORLD)

MPI_Gather(tmpd, NUMROWS, MPI_FLOAT, d, NUMROWS, MPI_FLOAT, 0, MPI_COMM_WORLD)
//NUMROWS = 2, 2 processes, 2x2 matrix

MPI Programming - Collectives

- MPI_Barrier – all processes wait at that line of code – a synchronization point

- MPI_Bcast – Broadcasting data from one process

- MPI_Scatter – Distribute data from master to all processes

- MPI_Gather – Collect data from all processes at master

- MPI_Allgather – Same as gather but all processes collect results

- MPI_Reduce – Aggregate results at master (recall `reduction` in OMP)

- MPI_Allreduce – Aggregate results at all processes

  refer https://computing.llnl.gov/tutorials/mpi/ for API details

MPI Programming – Point-to-Point

- In most MPI programs, communication is between a pair of processors.
  - Think other types of communication: Broadcast (one-to-all), Reduce(All-to-one), Scatter (one-to-several), Gather(several to one), All-to-All

- When is `send/receive` complete?
  - Synchronous / Asynchronous
  - Blocking / non-blocking
  - Buffered

MPI Programming – Point-to-Point

- Synchronous vs. Asynchronous
  - Synchronous: sender notified when message is received
  - Asynchronous: sender only knows that the message is sent
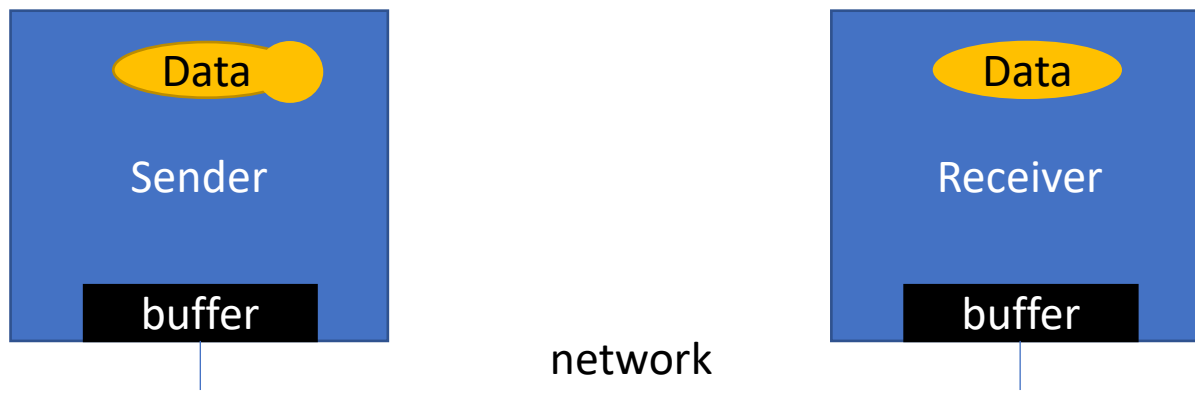
MPI Programming – Point-to-Point

- Blocking vs. Non-blocking
  - Blocking:
    - Sender waits until message is transmitted – buffer is empty
    - Receiver waits until message is received – buffer is full

  - Non-blocking
    - sender continues execution immediately after calling send